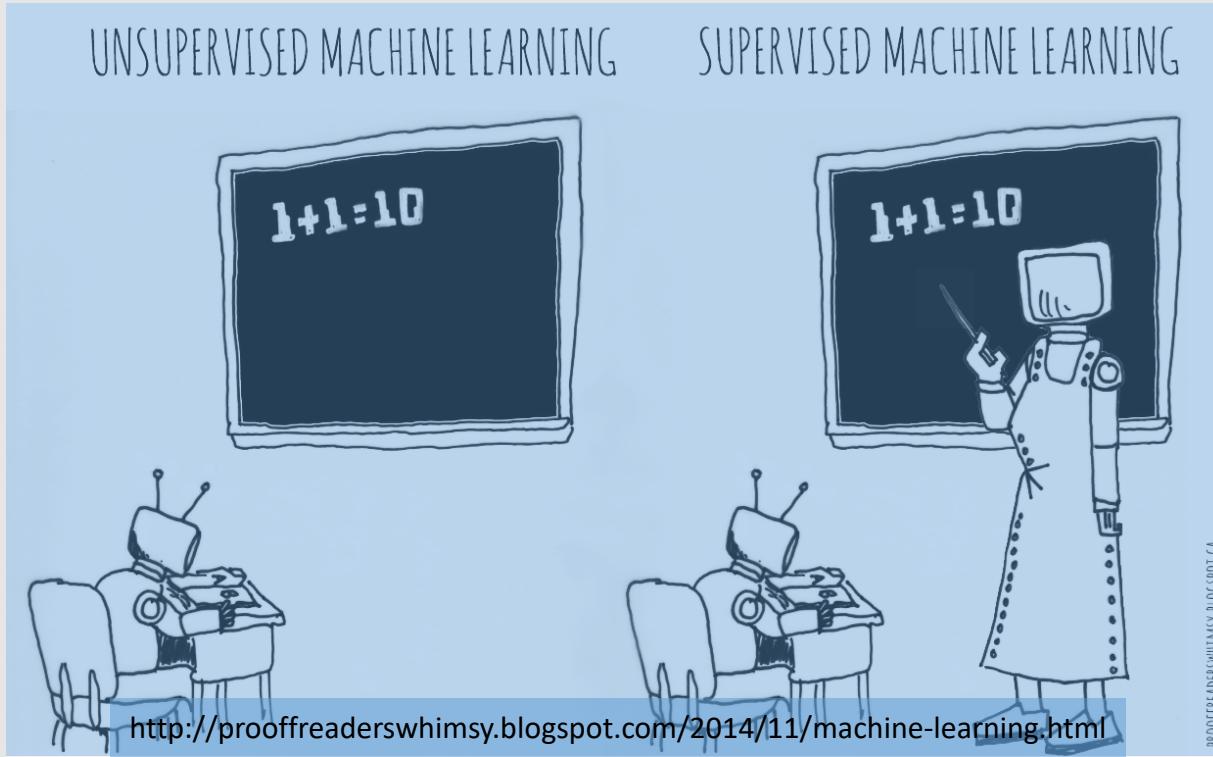


# Lecture 8: Unsupervised Learning



Haiping Lu

YouTube Playlist: <https://www.youtube.com/c/HaipingLu/playlists>

COM4059/6059: MLAI21@The University of Sheffield

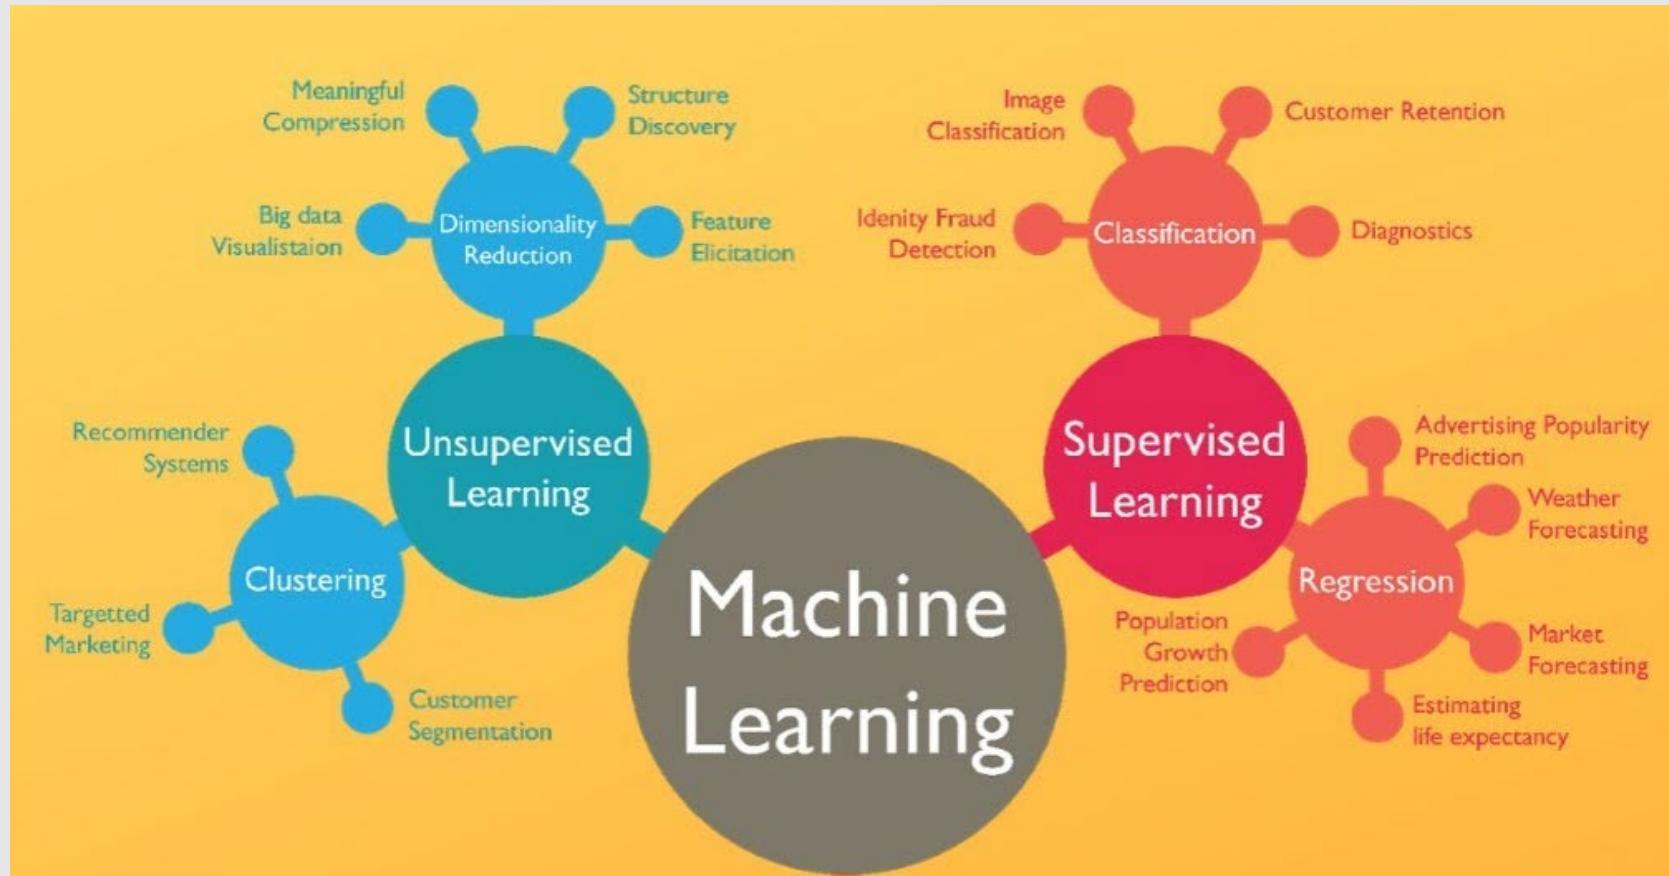
# Week 8 Contents / Objectives

- Why Unsupervised Learning?
- Principal Component Analysis (PCA)
- PCA Unboxing
- Clustering: from  $k$ -means to spectral
- Autoencoder

# Week 8 Contents / Objectives

- **Why Unsupervised Learning?**
- Principal Component Analysis (PCA)
- PCA Unboxing
- Clustering: from  $k$ -means to spectral
- Autoencoder

# Supervised vs Unsupervised



<https://i.morioh.com/2020/04/14/ff897f322fed.jpg>

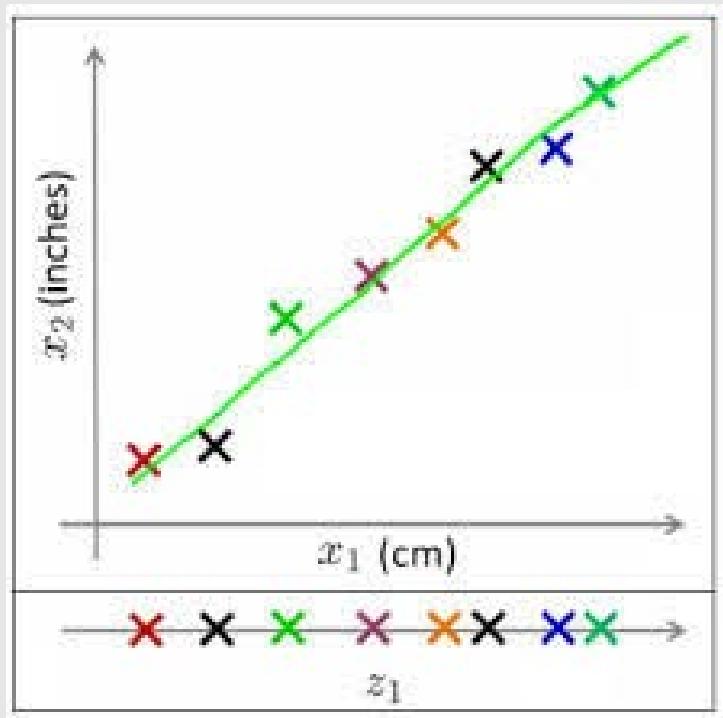
# Unsupervised Learning

- Supervised learning: each data point has a label
- Unsupervised learning: no labels for the data
  - Dimensionality reduction
  - Clustering

Machine Learning	Supervised	Unsupervised
Discrete output	Classification	Clustering
Continuous output	Regression	Dimensionality Reduction

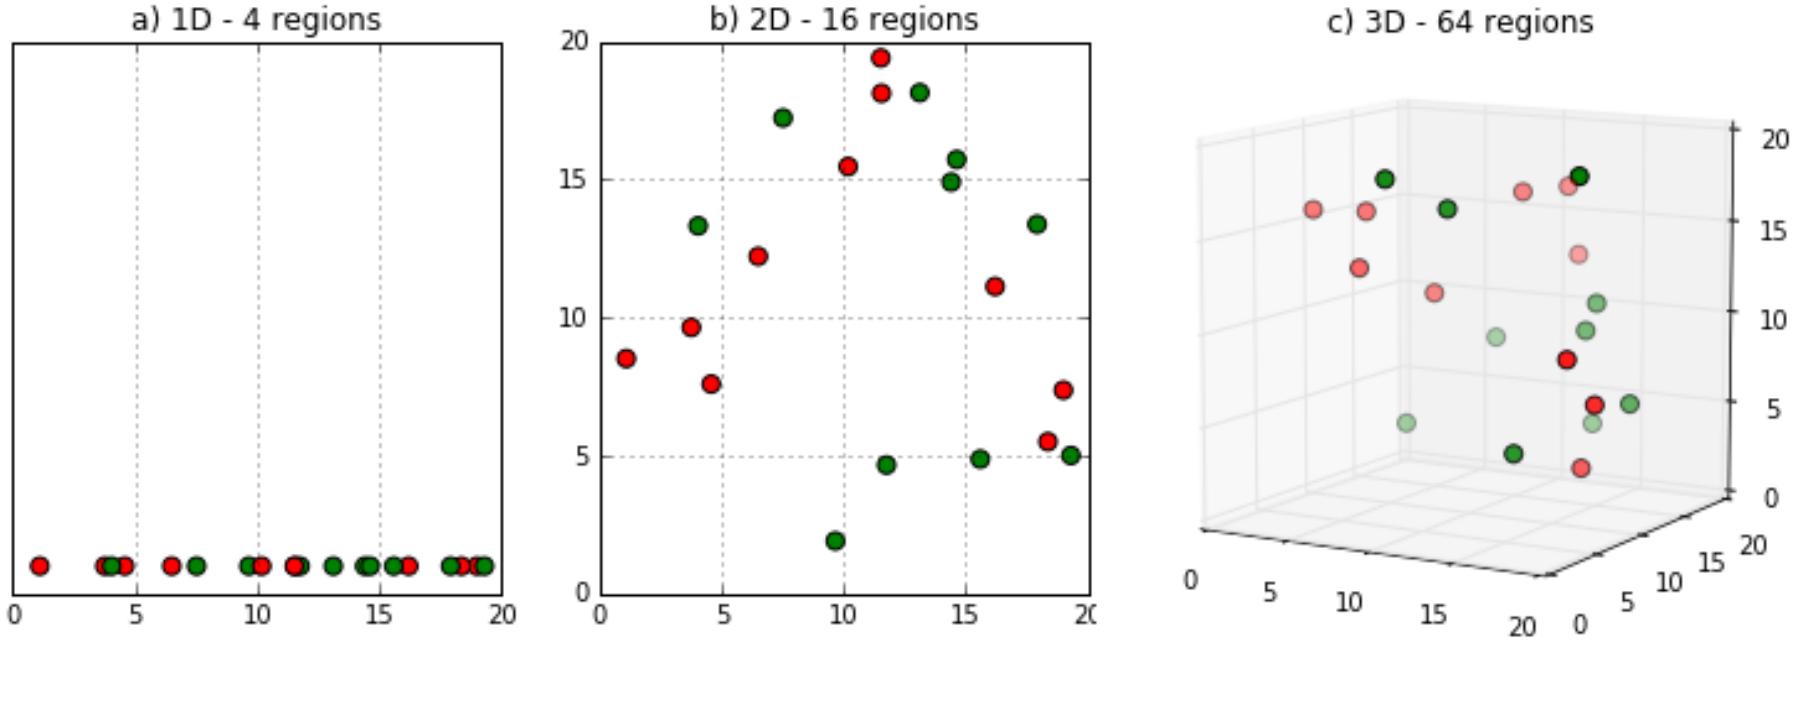
# Dimensionality Reduction (DR)

- $2 \rightarrow 1$



High Dimensional      Low Dimensional

Original data	Transformed
(1, 1.2)	1.15
(2, 2)	2
(3, 3.3)	3.1
...	...



# Why DR?

High D → Low D

- Curse of dimensionality
  - Nearest neighbours
- Reduce redundancy (correlation)
- Visualisation

# Question

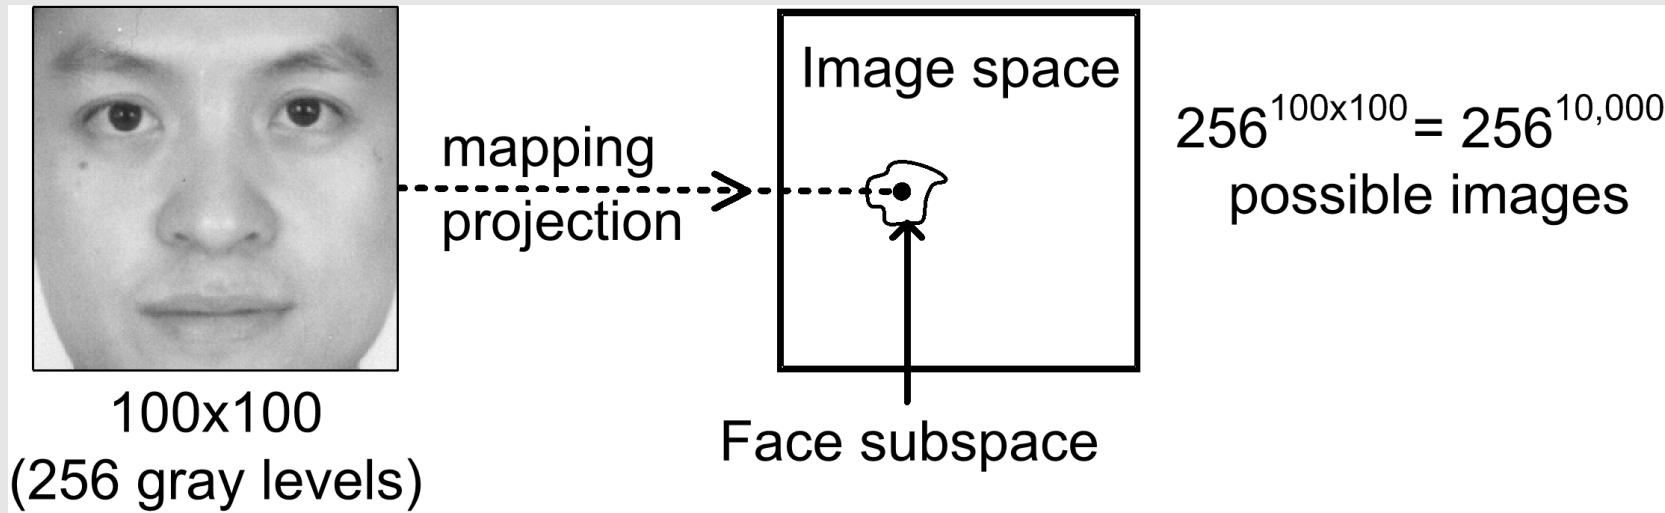
- USPS dataset handwritten digit
- Size: 64 x 57; binary (1-bit, BW)
- The binary image space contains much more than just this digit.



**How many possible images of this size and bit depth?**



$$2^{64 \times 57} = 2^{3648} = ?$$

# How About a Face?



# Low-D Subspace/Manifolds

- For high dimensional data with **structure**:
  - Fewer variations than dimensions
  - Data to live on a lower dimensional manifold
  - → Deal with them by looking for a lower dimensional embedding (or projection, transformation)



# Week 8 Contents / Objectives

- Why Unsupervised Learning?
- **Principal Component Analysis (PCA)**
- PCA Unboxing
- Clustering: from  $k$ -means to spectral
- Autoencoder

PCA?


Visualisation demo



# Principal Component Analysis

The idea:

1. Rotate the data with some rotation matrix  $R$  (change of basis) so that the new features are **uncorrelated**
2. Keep the dimension with the highest **variance** for DR



# Principal Component Analysis

- PCA (@Hotelling:analysis33): a linear embedding
- Rotate to find *directions* in data with **maximum variance**
- How do we find these directions?
  - Diagonalize the **sample covariance (scatter) matrix** of  $N$  samples  $\{\mathbf{x}^{(i)}, i = 1, \dots, N\}$

$$\mathbf{S} = \frac{1}{N} \sum_{i=1}^N (\mathbf{x}^{(i)} - \boldsymbol{\mu}) (\mathbf{x}^{(i)} - \boldsymbol{\mu})^\top$$

# Principal Component Analysis

- Given data  $\{\mathbf{x}^{(i)}\}$ , PCA finds **orthogonal** directions defined by a projection (rotation)  $\mathbf{U}$  capturing the **maximum variance** in the data
- Solution: eigenvectors of  $\mathbf{S} \rightarrow \mathbf{U}$
- PCA representation  $\mathbf{y} = \mathbf{U}^\top \mathbf{x}$
- **Question:** Given the PCA representation  $\mathbf{y}$ , how to obtain an approximation/reconstruction of  $\mathbf{x}$ ?

$$\hat{\mathbf{x}} = \mathbf{U}\mathbf{y}$$

# Representation & Reconstruction

- Face  $\mathbf{x}$  in  $k$  “face space” coordinates



$$\begin{aligned}\mathbf{x} &\rightarrow [\mathbf{u}_1^\top (\mathbf{x} - \boldsymbol{\mu}), \dots, \mathbf{u}_k^\top (\mathbf{x} - \boldsymbol{\mu})] \\ &= [w_1, \dots, w_k]\end{aligned}$$

- Reconstruction: eigenvectors as orthonormal basis vectors

$$\begin{matrix} \text{[Face Image]} & = & \text{[Mean Face]} & + & \text{[Eigenfaces (7 columns)]} \end{matrix}$$

$$\hat{\mathbf{x}} = \boldsymbol{\mu} + w_1 \mathbf{u}_1 + w_2 \mathbf{u}_2 + w_3 \mathbf{u}_3 + w_4 \mathbf{u}_4 + \dots$$

# Reconstruction


$k = 4$



$k = 200$



$k = 400$



After computing eigenfaces using 400 face images from the ORL face database

# PCA Ingredients

- **Data:** +pre-processing, e.g.,  $\mathcal{N}(0,1)$
- **Model**
  - Structure/Architecture: linear projection  $y = U^T x$
  - **Hyper-parameter:** lower dimension  $k$
  - Parameters (theta): the principal components (eigenvectors)
- Evaluation metric: max variance
- Optimisation: eigen-decomposition

# Week 8 Contents / Objectives

- Why Unsupervised Learning?
- Principal Component Analysis (PCA)
- **PCA Unboxing**
- Clustering: from  $k$ -means to spectral
- Autoencoder

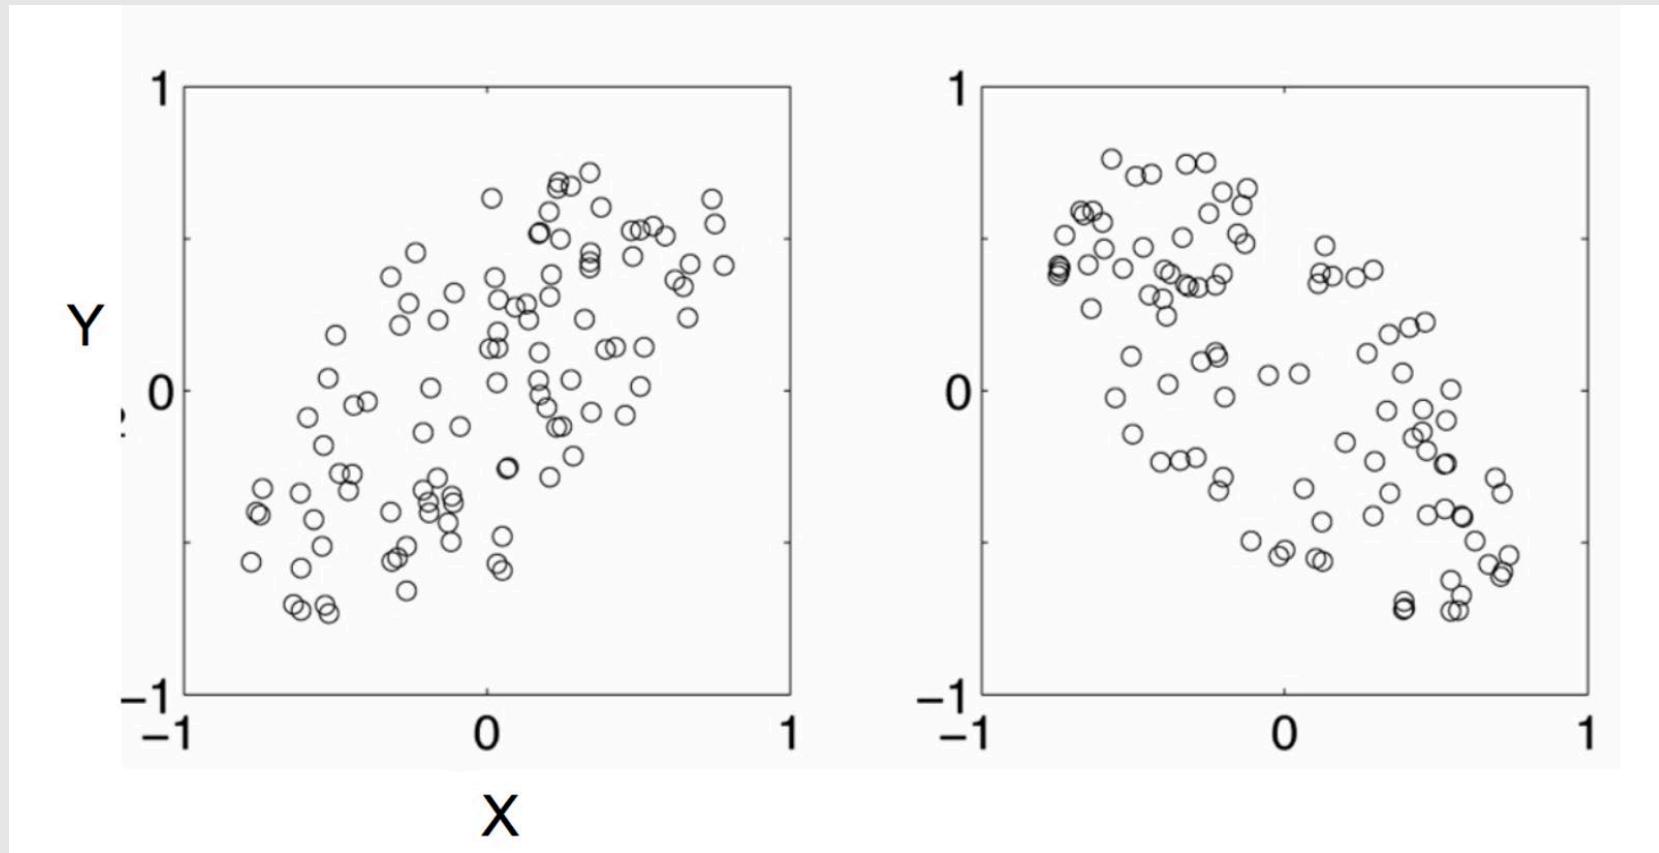
# Variance → Covariance (Matrix)

- Variance and covariance:
  - Measure of the “spread” of a set of points around their *center of mass* (mean)
- Variance (scalar): *Scalar and matrix*
  - Measure of the deviation from the mean for points in **one dimension**
- Covariance (matrix):
  - Measure of how much each of the dimensions vary from the mean with **respect to each other**



- Covariance is measured between two dimensions
- Covariance sees if there is a **relation** between the two dimensions
- Covariance between one dimension is the variance (degeneration)

# Positive/Negative Covariance



Positive: Both dimensions increase or decrease together

Negative: While one increase the other decrease

# Covariance *Matrixp*

- Find relationships between dimensions in high dimensional data sets  $\mathbf{X}$
- Scatter matrix  $\mathbf{S}$ : sample-based estimation of covariance matrix ( $i$ : sample;  $j/k$ : variable)

$$s_{jk} = \frac{1}{N} \sum_{i=1}^N (X_{ij} - E(X_j))(X_{ik} - E(X_k))$$


The Sample mean

- Uncorrelated variables  $\rightarrow$  covariance = 0
- Diagonal covariance mat  $\rightarrow$  all variables are uncorrelated

# PCA Derivation – Max Variance

*keep the information*

- Scatter mat for the input  $\mathbf{S} = \frac{1}{N} \sum_{i=1}^N (\mathbf{x}^{(i)} - \boldsymbol{\mu}) (\mathbf{x}^{(i)} - \boldsymbol{\mu})^\top$
- **Question:** what is the scatter mat for the projections?

$$\mathbf{U}^\top \mathbf{S} \mathbf{U}$$

- First PC: **maximise the variance in the projected space**, i.e. the variance of  $y = \underline{\mathbf{u}_1^\top \mathbf{x}}$

$$\begin{aligned}\text{var}(y) &= \frac{1}{N} \sum_{i=1}^N (y^{(i)} - \mu_y)^2 \\ &= \frac{1}{N} \sum_{i=1}^N \mathbf{u}_1^\top (\mathbf{x}^{(i)} - \boldsymbol{\mu}_{\mathbf{x}}) (\mathbf{x}^{(i)} - \boldsymbol{\mu}_{\mathbf{x}})^\top \mathbf{u}_1 \\ &= \mathbf{u}_1^\top \mathbf{S} \mathbf{u}_1\end{aligned}$$

# PCA: Max Variance → Eigenvalue

- Find the first direction  $\mathbf{u}_1$  via a unit-norm constrained optimisation, using Lagrange multipliers:

$$L(\mathbf{u}_1, \lambda_1) = \mathbf{u}_1^\top \mathbf{S} \mathbf{u}_1 + \lambda_1 (1 - \mathbf{u}_1^\top \mathbf{u}_1)$$

- Gradient w.r.t.  $\mathbf{u}_1$ :  $\frac{dL(\mathbf{u}_1, \lambda_1)}{d\mathbf{u}_1} = 2\mathbf{S}\mathbf{u}_1 - 2\lambda_1\mathbf{u}_1$

- Set to 0 and rearrange:  $\mathbf{S}\mathbf{u}_1 = \lambda_1\mathbf{u}_1 \rightarrow$  First PC

- Eigenvalue problem

- **Question:** many solutions (eigen-pairs), which to choose?

$$\text{var}(y) = \mathbf{u}^\top \mathbf{S} \mathbf{u} = \lambda \mathbf{u}^\top \mathbf{u} = \lambda$$

# PCA Solution

- Further directions: **orthogonal** (uncorrelated) to the first and each others → top  $k$  eigenvectors of  $\mathbf{S}$ 
  - Eigenvectors: basis functions, principal components
  - Eigenvalue: the **variance** captured respectively
- **Questions**
  - For  $\mathbf{u}_1$ , what if we do not have the unit-norm constraint?
$$L(\mathbf{u}_1, \lambda_1) = \mathbf{u}_1^\top \mathbf{S} \mathbf{u}_1 + \cancel{\lambda_1 (1 - \mathbf{u}_1^\top \mathbf{u}_1)}$$
→ Trivial solution: infinity
  - For further directions: what if we do not require them to be orthogonal?  
→ We will have the same  $\mathbf{u}_1$ , useless solution

# PCA: Max Variance $\leftrightarrow$ Min MSE

*Principle* Consider the first PC with the projection vector  $\mathbf{u}$ . We assume *zero-mean*, i.e. *centered* data

**Maximum Variance Direction:** 1<sup>st</sup> PC = a vector  $\mathbf{u}$  such that projection on it captures max variance in the data

$$\frac{1}{N} \sum_{i=1}^N (\mathbf{u}^T \mathbf{x}^{(i)})^2 = \mathbf{u}^T \mathbf{X} \mathbf{X}^T \mathbf{u}$$

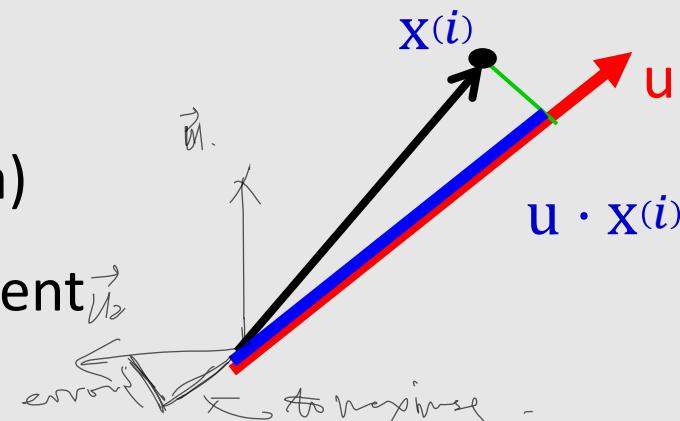
**Minimum Reconstruction Error:** 1<sup>st</sup> PC = a vector  $\mathbf{u}$  such that projection on it yields minimum MSE reconstruction

$$\frac{1}{N} \sum_{i=1}^N \|\mathbf{x}^{(i)} - (\mathbf{u}^T \mathbf{x}^{(i)}) \mathbf{u}\|^2$$

*error.*  
blue<sup>2</sup> + green<sup>2</sup> = black<sup>2</sup>

black<sup>2</sup> is fixed (it's just the data)


So, maximizing blue<sup>2</sup> is equivalent  
to minimizing green<sup>2</sup>



# How many ( $k$ ) PCs to keep?

*principal components*

- Pick based on percentage of variance captured / lost
  - Variance captured: the variance of the projected data
  - Pick smallest  $k$  that explains a certain percentage of variance
$$(\text{Sum of first } k \text{ EVs}) / (\text{sum of all EVs})$$
- Look for an “elbow” in scree plot (plot of explained variance or eigenvalues)



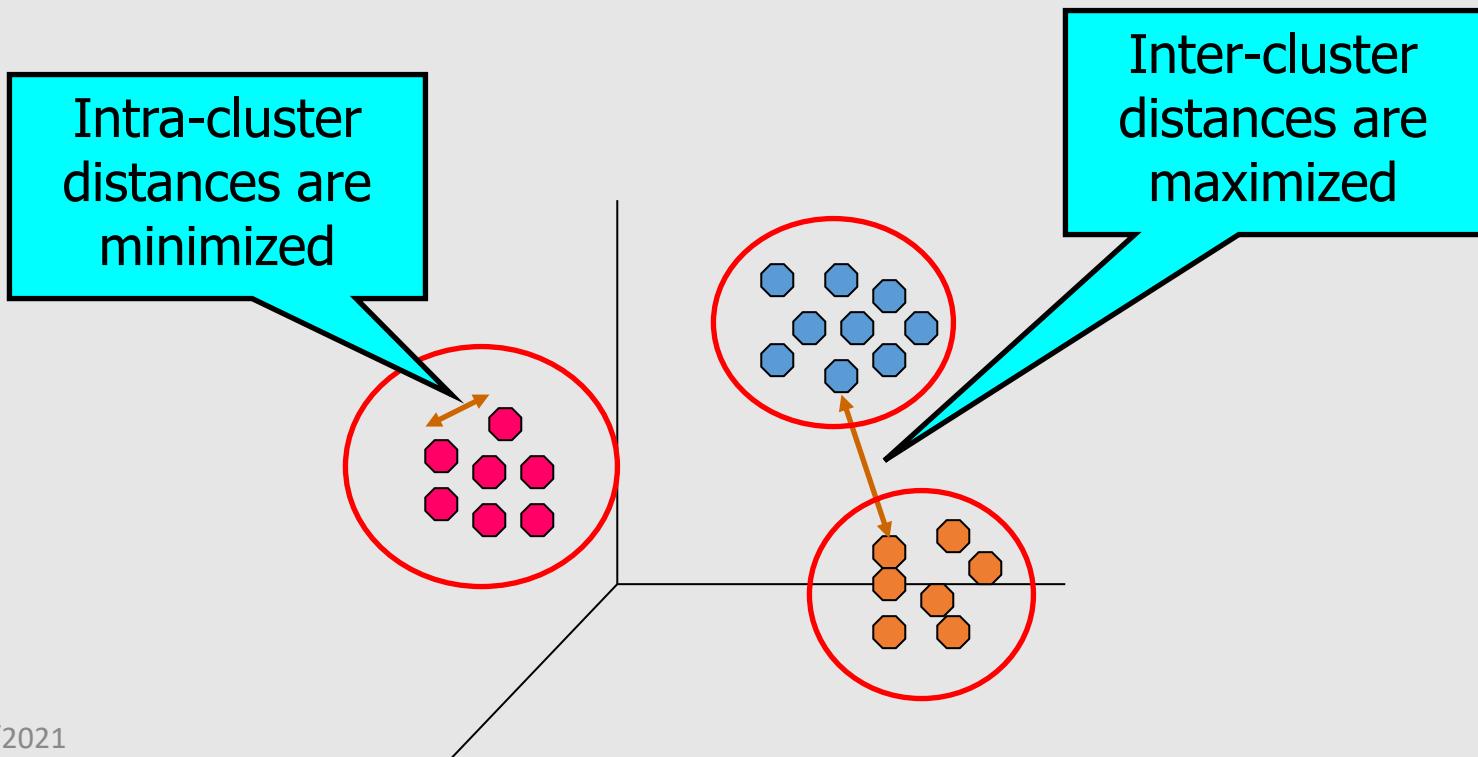
# Week 8 Contents / Objectives

- Why Unsupervised Learning?
- Principal Component Analysis (PCA)
- PCA Unboxing
- **Clustering: from  $k$ -means to spectral**
- Autoencoder

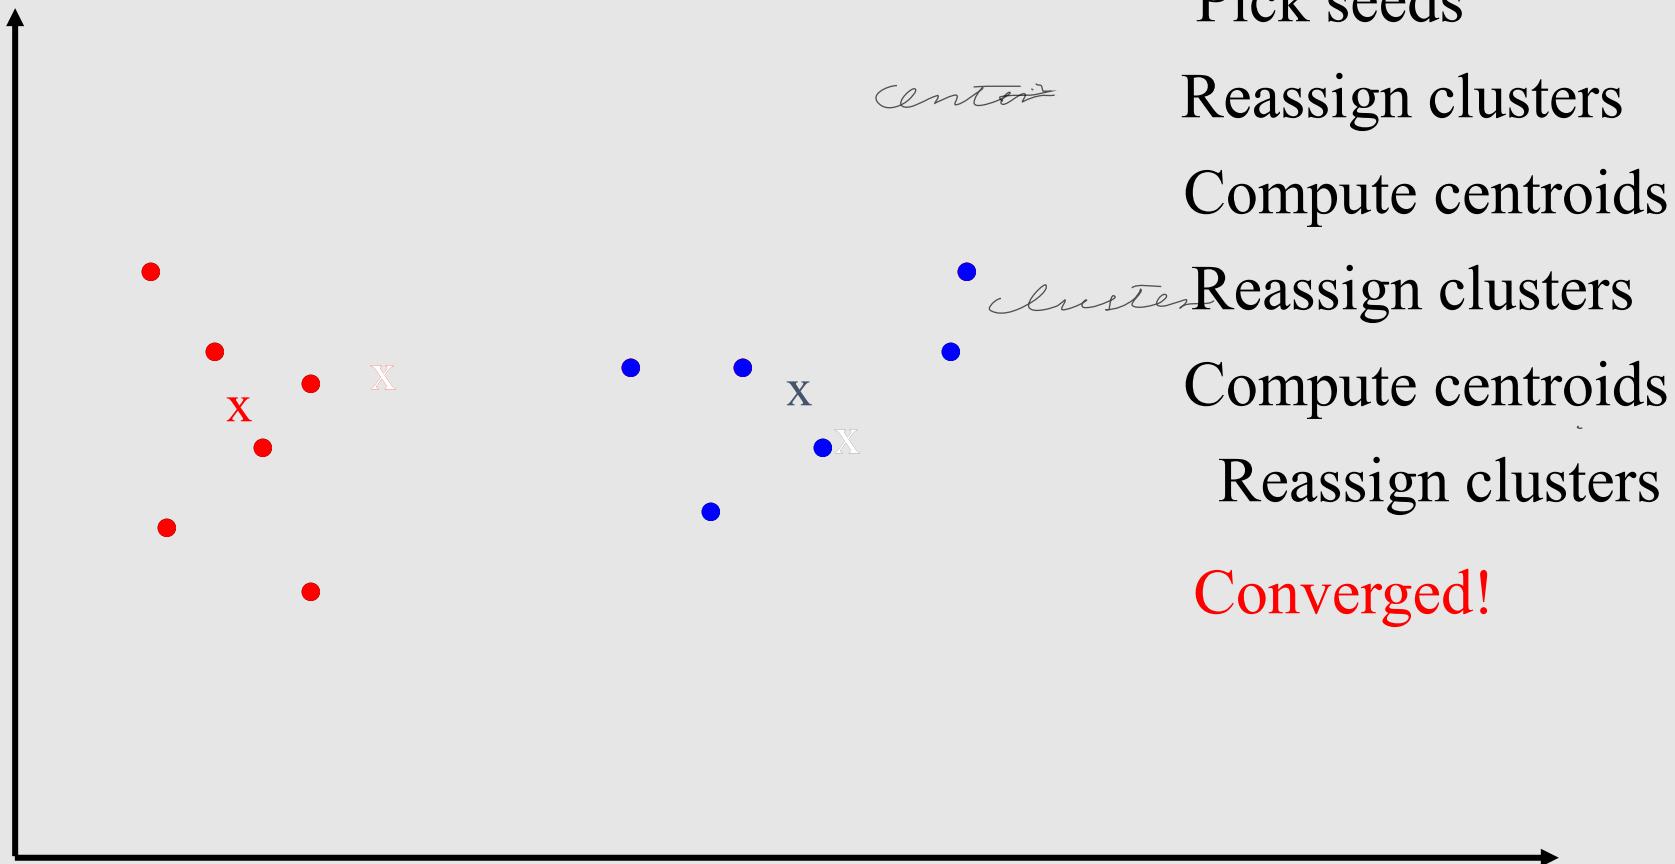
# What is Clustering?

*Defines a metric to separate clusters/boundaries*

- Find grouping of objects such that the objects in a group will be similar or more related to one another and those in different groups will be different or less related



# $k$ -means Example ( $k=2$ )



# *k*-means & PCA

*They will create  
labels based on algorithm*

Machine Learning	Supervised	Unsupervised
Discrete output	Classification	<i>Labels included into clustering</i>
Continuous output	Regression	Dimensionality Reduction

- The objective of *k*-means clustering: to minimize the within-cluster scatter

$$\min \sum_{j=1}^k \sum_{i \text{ allocated to } j} \left( \mathbf{x}^{(i)} - \boldsymbol{\mu}^{(j)} \right)^T \left( \mathbf{x}^{(i)} - \boldsymbol{\mu}^{(j)} \right)$$

*j - center to minimize variance ↘*

- Similarity with PCA: also measure of the “spread” of a set of points around their *center of mass* (mean)
- Clustering analogy
  - Classification w/o labelled training data
  - Extreme dimensionality reduction (to a **cluster label**)


# $k$ -means Clustering Ingredients

- **Data:** +pre-processing, e.g.,  $\mathcal{N}(0,1)$
- **Model**
  - Structure/Architecture: linear separation btw clusters
  - **Hyper-parameter:** #clusters  $k$
  - **Parameters (theta):** the  $k$  cluster centroids  
*to learn which parameter will separate the clusters with the maximum variance*
- Evaluation metric: within-cluster scatter
- Optimisation: expectation maximisation (EM), kind of gradient descent


# $k$ -means Has a Problem

- $k$ -means: centroid-based, for compactness (scatter)
- Spectral clustering: **connectivity**

$k$ -means





Spectral clustering



[Shi & Malik '00; Ng, Jordan, Weiss NeurIPS '01]

# Spectral Clustering

- Group points based on **links** in a **graph**



- Image as graph: segmentation as clustering *connectedivity*



# How to Create the Graph?

- Gaussian kernel → compute similarity between objects

$$\mathbf{W}(i, j) = \exp \frac{-|x_i - x_j|^2}{\sigma^2}$$

- One could create


- A fully connected graph (~ FC layer)

*the most similar e.g. social network*

- K-nearest neighbour graph: each node is only connected to its K-nearest neighbours (~ convolutional layer, local connectivity)

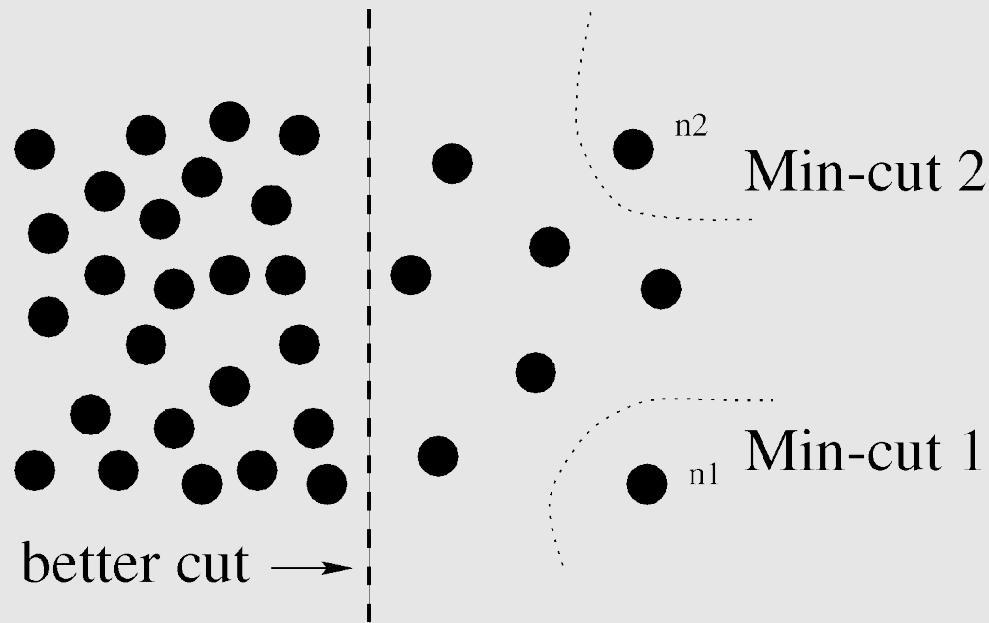
- Clustering → Graph cut/partition

- Objective: **minimise** cut



# Min Cut = Good Cut?

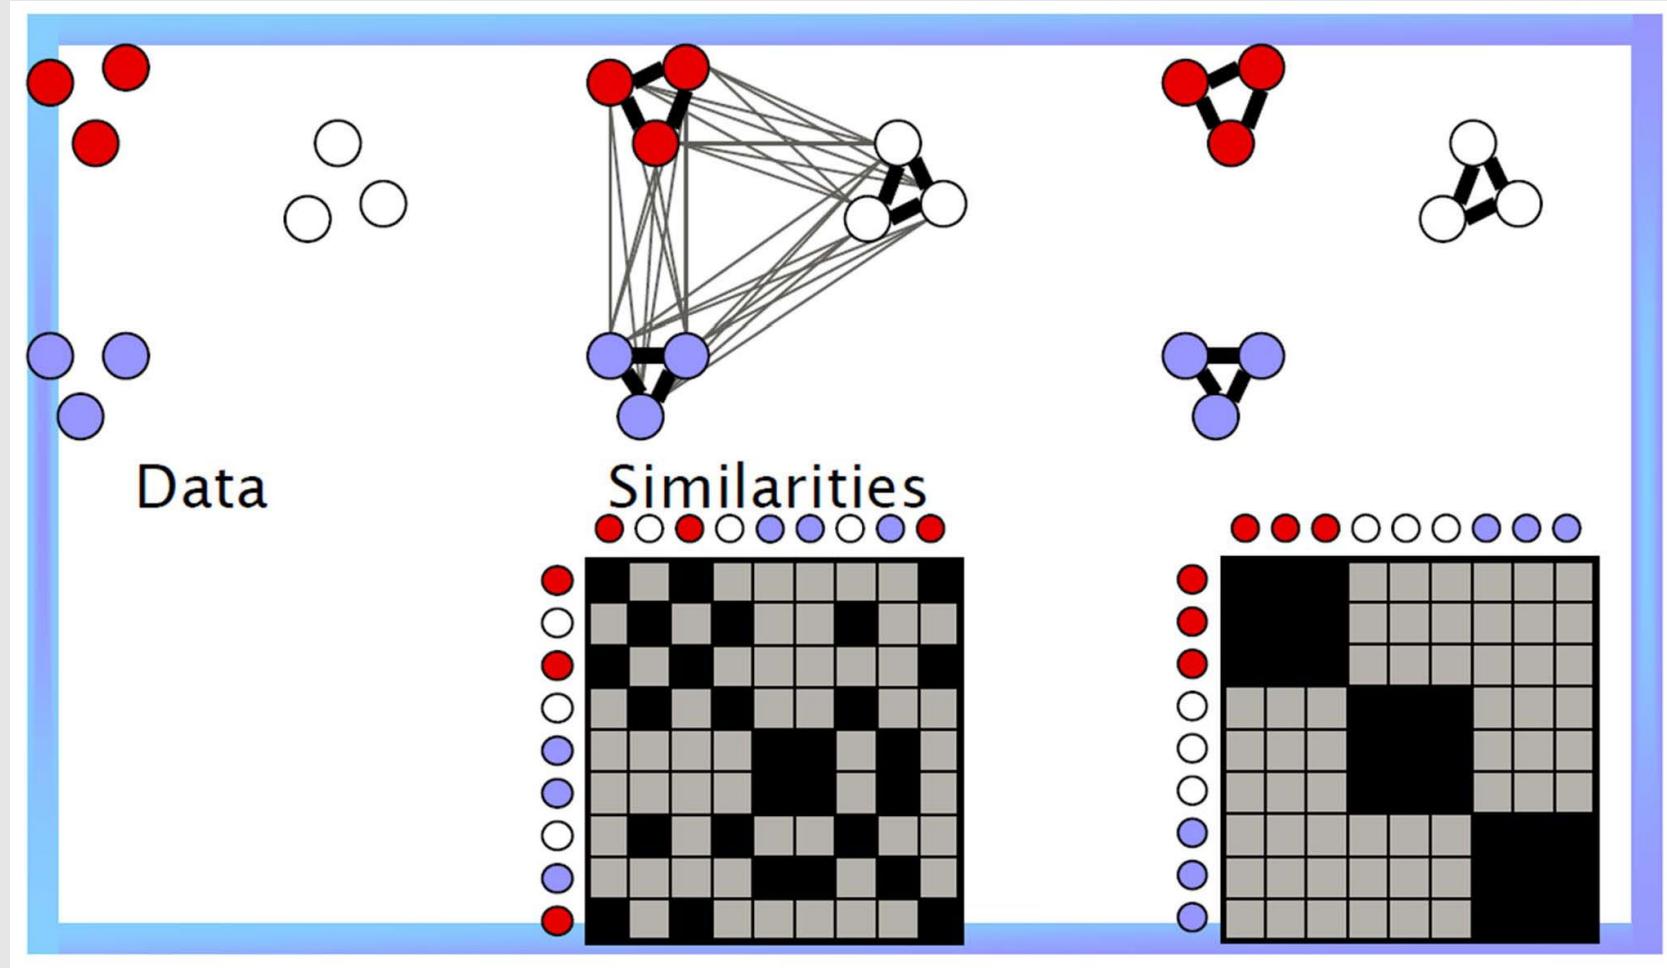
- A case where minimum cut gives a bad partition



- Solution: **Normalise** the cut

[Shi & Malik '00]


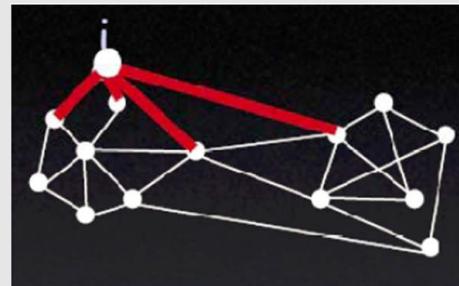
# Graph Clustering Process



# Graph Terminologies



- Degree of nodes

$$d_i = \sum_j w_{i,j}$$




- Volume of a set

$$vol(A) = \sum_{i \in A} d_i, A \subseteq V$$



# Graph Cut

- Consider a partition of the graph into two parts A & B



**Question**

$$cut(A, B) =$$

- $cut(A, B)$ : sum of the weights of the set of edges that connect the two groups

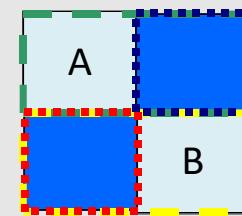
$$cut(A, B) = \sum_{i \in A, j \in B} w_{ij}$$

- Goal: find the partition that minimizes the cut

# Normalized Cut (Ncut)

- Consider the **connectivity** between groups **relative** to the volume of each group

$$Ncut(A, B) = \frac{cut(A, B)}{Vol(A)} + \frac{cut(A, B)}{Vol(B)}$$



$$Ncut(A, B) = cut(A, B) \frac{Vol(A) + Vol(B)}{Vol(A)Vol(B)}$$

# Solving/Minimising Ncut

- Compute the similarity matrix  $\mathbf{W}$  :  $\mathbf{W}(i, j) = w_{i,j}$
- Compute the degree matrix  $\mathbf{D}$  :  $\mathbf{D}(i, i) = \sum_j w_{i,j}$
- Solve a generalised **eigenvalue** problem (relaxed Ncut)

$$\min_{\mathbf{y}} \mathbf{y}^\top (\mathbf{D} - \mathbf{W}) \mathbf{y} \text{ s.t. } \mathbf{y}^\top \mathbf{D} \mathbf{y} = 1 \Rightarrow (\mathbf{D} - \mathbf{W}) \mathbf{y} = \lambda \mathbf{D} \mathbf{y}$$

$(\mathbf{D} - \mathbf{W})$  : Laplacian matrix

- Solution:
  - Bipartition: use the eigenvector with the second smallest eigenvalue to partition the graph into two parts
    - Splitting point: minimum Ncut (plot). *all the points in the graph.*
  - K-way partition: **k-means clustering** of multiple eigenvectors
    - Graph embedding (dimensionality reduction) → eigenvectors

# Recap: Power of Transform

- Logistic regression **transforms** classification into linear regression of the log odds, modelling the probability of the predicted output rather than the output itself.
- Spectral clustering **transforms** non-linear clustering into (linear)  $k$ -means clustering of generalised eigenvectors based on the similarity graph, modelling the connectivity of data points rather than themselves.



# Spectral Clustering Ingredients

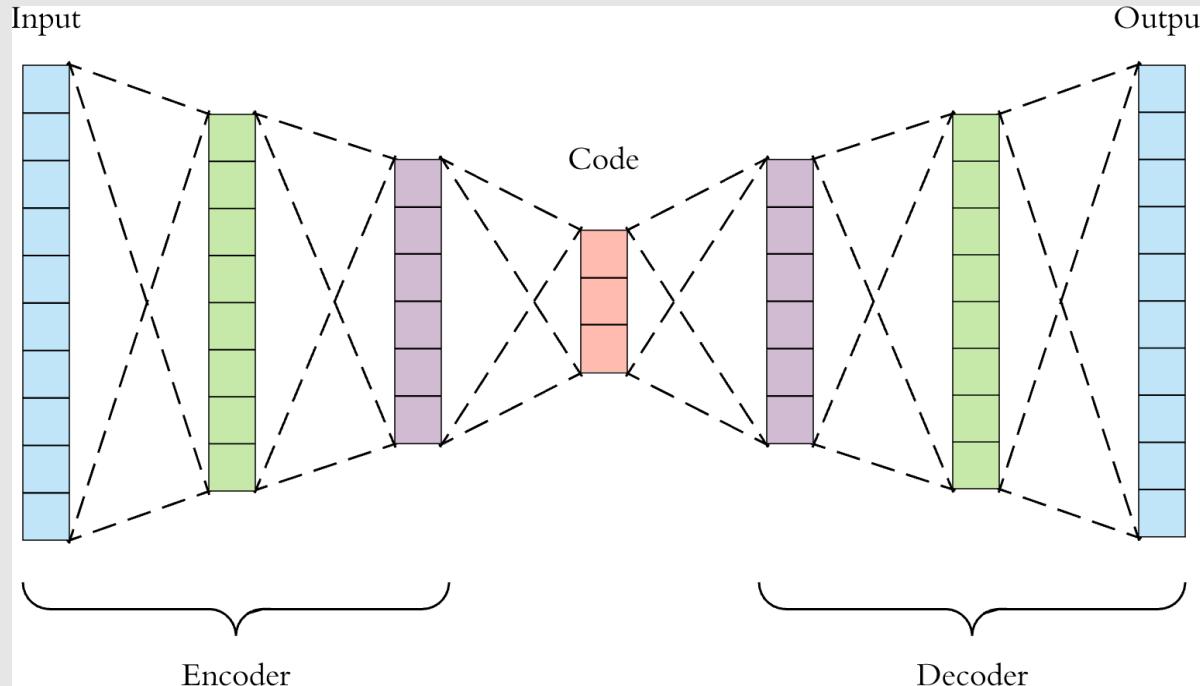
- **Data:** +pre-processing, e.g.,  $\mathcal{N}(0,1)$
- **Model**
  - Structure/Architecture: spectral (nonlinear) separation btw clusters
  - **Hyper-parameter:** #clusters  $k$  (+#eigenvectors), kernel bandwidth  $\sigma$
  - Parameters (theta): the (generalised) eigenvectors
- Evaluation metric: normalised (graph) cut
- Optimisation: eigen-decomposition (and expectation maximisation in  $k$ -means)

# Week 8 Contents / Objectives

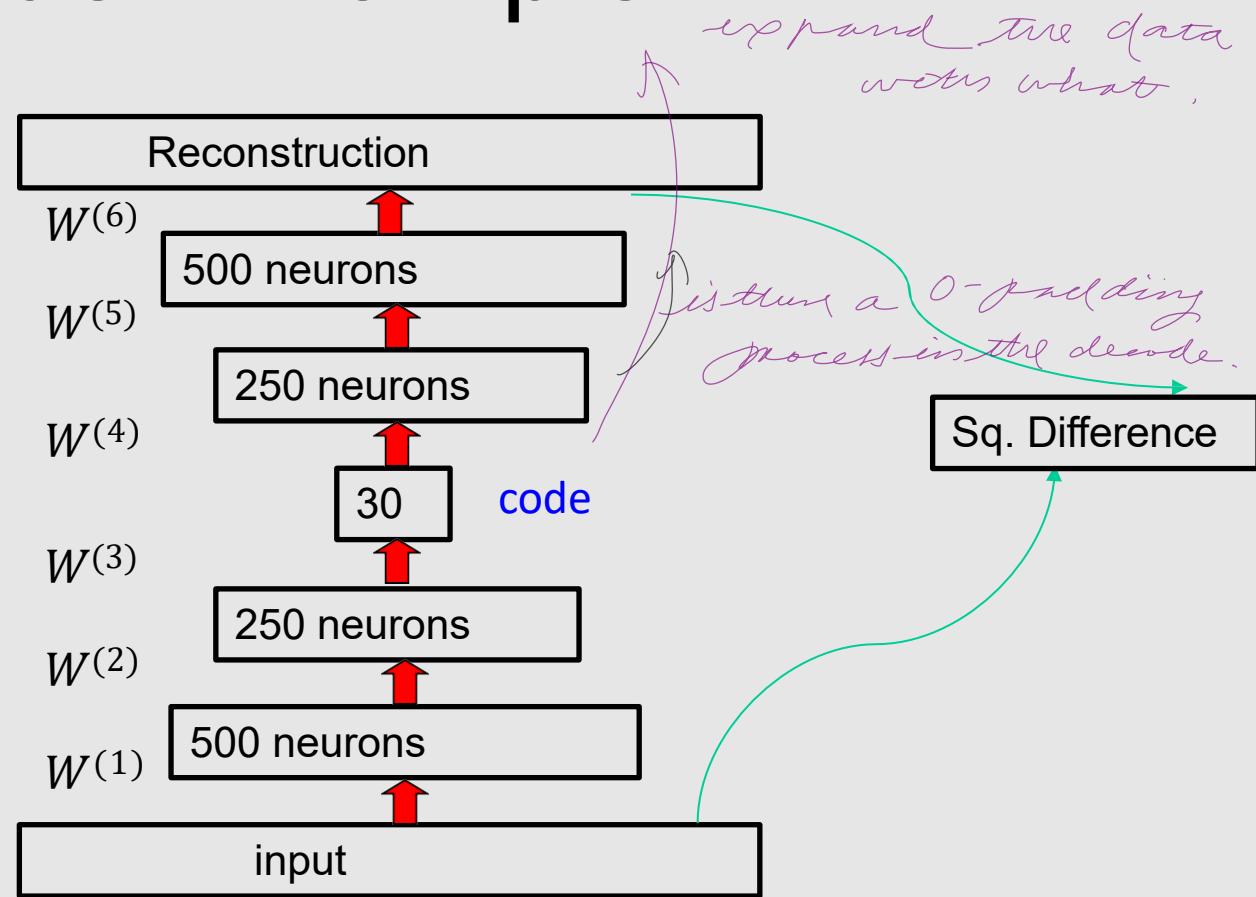
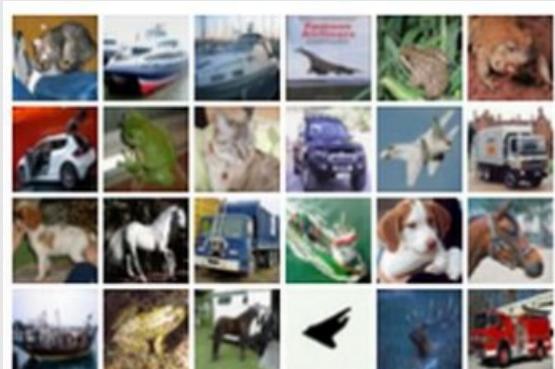
- Why Unsupervised Learning?
- Principal Component Analysis (PCA)
- PCA Unboxing
- Clustering: from  $k$ -means to spectral
- **Autoencoder**

# Autoencoder

- **Encoder:** compress data or extract features
- **Decoder:** generate images given a new code
- **Bottleneck (code):** to make it non-trivial, much smaller dimension as the latent representation



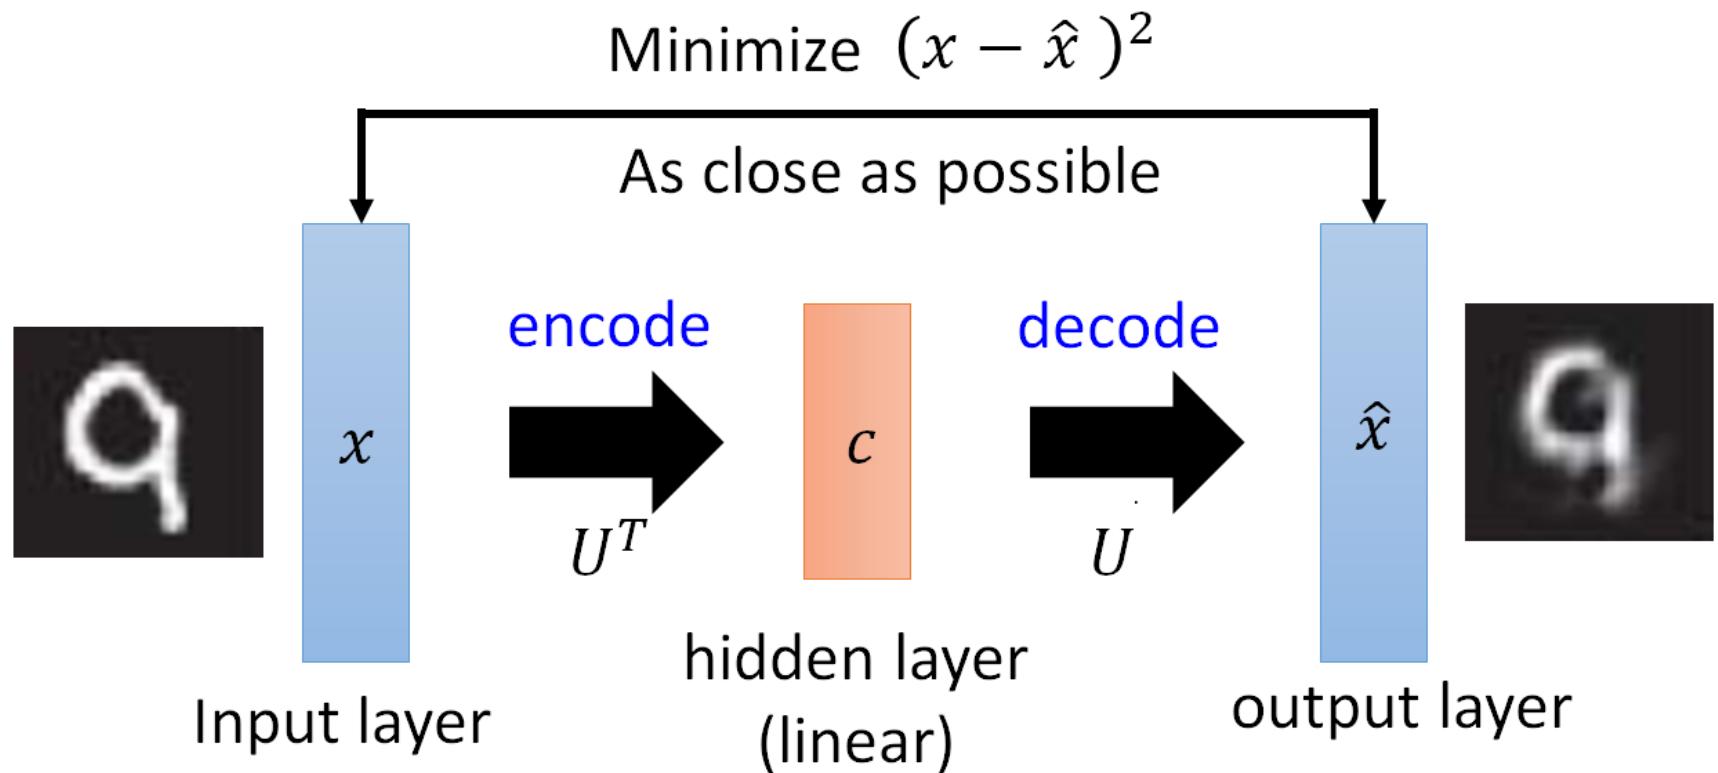
# Autoencoder Example



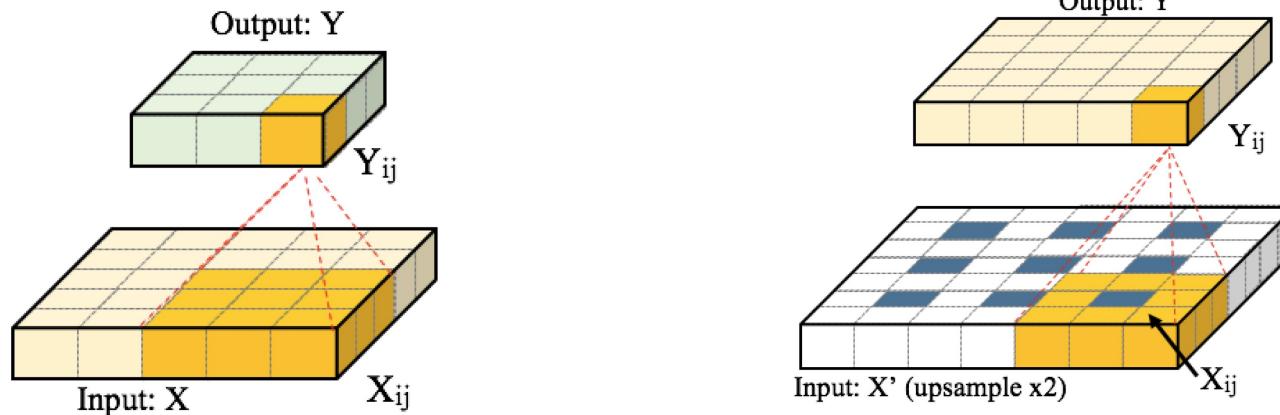
- Find the weights minimising the difference between the input and reconstruction
- The **code** layer (bottleneck) is a low-dimensional summary of the input

# PCA as Linear Autoencoder

- PCA: autoencoder w/t single-layer encoder/decoder
- Weight sharing between the encoder and decoder



# Transpose Convolution Layer



(a) Convolutional layer: the input size is  $W_1 = H_1 = 5$ ; the receptive field  $F = 3$ ; the convolution is performed with stride  $S = 1$  and no padding ( $P = 0$ ). The output  $Y$  is of size  $W_2 = H_2 = 3$ .

(b) Transposed convolutional layer: input size  $W_1 = H_1 = 3$ ; transposed convolution with stride  $S = 2$ ; padding with  $P = 1$ ; and a receptive field of  $F = 3$ . The output  $Y$  is of size  $W_2 = H_2 = 5$ .

<https://www.mdpi.com/2072-4292/9/6/522/htm>

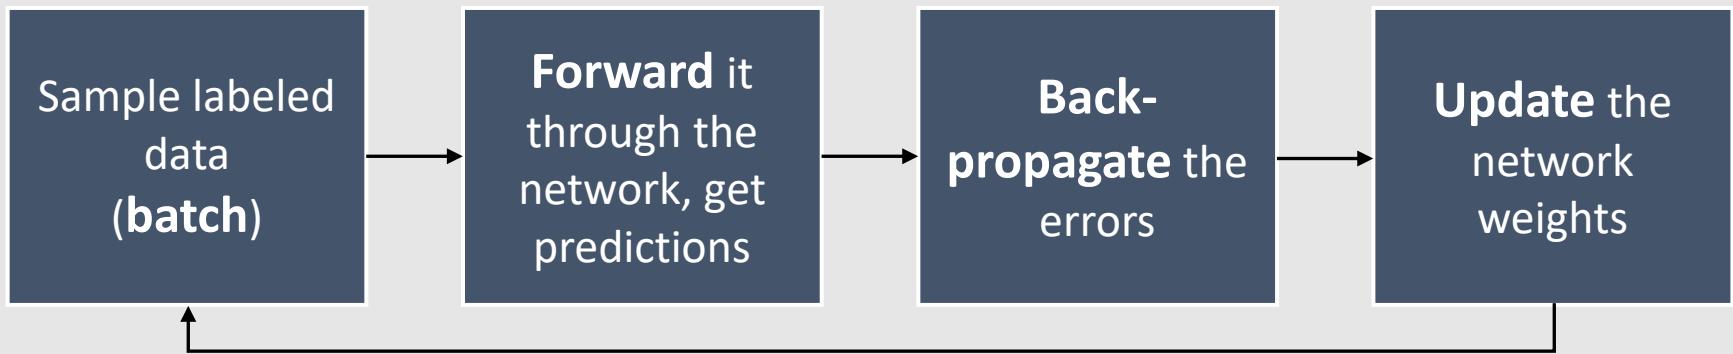
More at [https://github.com/vdumoulin/conv\\_arithmetic](https://github.com/vdumoulin/conv_arithmetic)

# Convolutional Autoencoder (Lab)

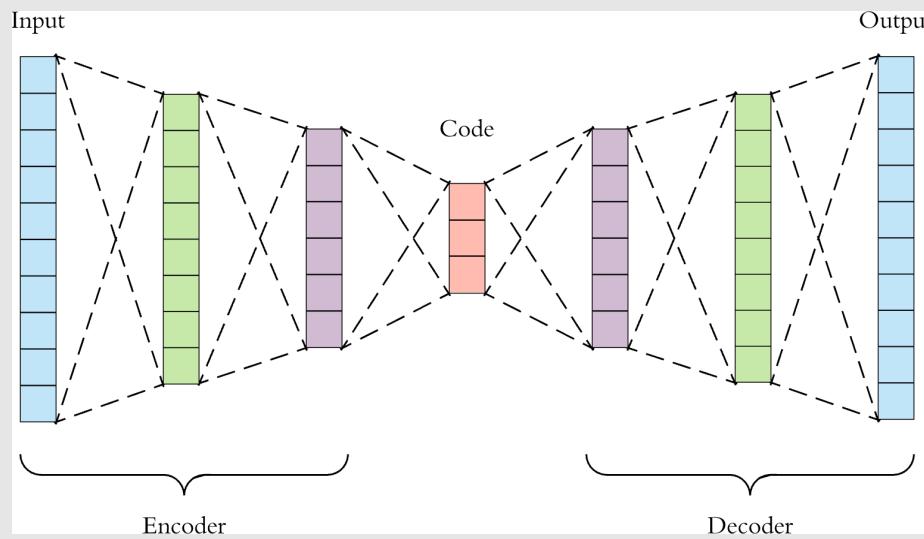
```
class Autoencoder(nn.Module):
    def __init__(self):
        super(Autoencoder, self).__init__()
        self.encoder = nn.Sequential(
            # 1 input image channel, 16 output channel, 3x3 square convolution
            nn.Conv2d(1, 16, 3, stride=2, padding=1),
            nn.ReLU(),
            nn.Conv2d(16, 32, 3, stride=2, padding=1),
            nn.ReLU(),
            nn.Conv2d(32, 64, 7)
        )
        self.decoder = nn.Sequential(
            nn.ConvTranspose2d(64, 32, 7),
            nn.ReLU(),
            nn.ConvTranspose2d(32, 16, 3, stride=2, padding=1, output_padding=1),
            nn.ReLU(),
            nn.ConvTranspose2d(16, 1, 3, stride=2, padding=1, output_padding=1),
            nn.Sigmoid() #to range [0, 1]
        )

    def forward(self, x):
        x = self.encoder(x)
        x = self.decoder(x)
        return x
```

# Training



Data → Model → Metric → Optimisation



# Autoencoder Ingredients

- **Data:** + pre-processing, e.g.,  $\mathcal{N}(0,1)$
- **Model**
  - Structure/Architecture: layers defined in nn.module
  - **Hyper-parameter:** layer specs, e.g. #layers #channels, kernel size
  - Parameters (theta): layer weights and biases
- Evaluation metric: MSE or other
- Optimisation: backprop, SGD or the like

## Acknowledgement

- The slides used materials from:  
*Lisa Zhang, Michael Guerzhoy,  
Neil Lawrence, Derek Hoiem,  
Pandu Nayak, Prabhakar  
Raghavan, Fereshteh Sadeghi,  
Aarti Singh, David Sontag, James  
Hays, Alan Fern, Tommi  
Jaakkola, Jure Leskovec*

## Recommended Reading

- The [PCA book](#) (from a UIUC link)
  - Chapter on PCA in most machine learning books
  - Chapter on clustering in most machine learning books
  - The [normalized cut paper](#) in 2000
  - Chapter on autoencoder in [the Deep Learning Book](#)
- 
- Wikipedia entries on covered topics
  - Scikit-learn/PyTorch documentations
  - The lab notebook and references